HIVE_MINI_PROJECT_1

Download Dataset 1 - https://drive.google.com/file/d/1WrG-9qv6atP-W3P_-gYln1hHyFKRKMHP/view

Download Dataset 2 - https://drive.google.com/file/d/1-JIPCZ34dyN6k9CqJa-Y8yxIGq6vTVXU/view

1. **Create a schema based on the given dataset**

```
[cloudera@quickstart ~]$ hdfs dfs -mkdir Agent_Details
[cloudera@quickstart ~]$ hdfs dfs -put /home/cloudera/Hive-Mini-Proj-1/AgentLogingReport.csv Agent_Details/

[cloudera@quickstart ~]$ hdfs dfs -put /home/cloudera/Hive-Mini-Proj-1/AgentPerformance.csv Agent_Details/
[cloudera@quickstart ~]$ hdfs dfs -ls Agent_Details
Found 2 items
-rw-r--r--   1 cloudera cloudera      55351 2023-03-21 03:18 Agent_Details/AgentLogingReport.csv
-rw-r--r--   1 cloudera cloudera     109853 2023-03-21 03:22 Agent_Details/AgentPerformance.csv
```

Hive> Create table AgentLogingReport (

sr_no int,

Agent string,

Date date,

Login string,

Logout string,

Duration string

)

row format delimited

fields terminated by ','

tblproperties ("skip.header.line.count" = "1");

Hive> Create table AgentPerformance

(

sr_no int,

Date date,

Agent_Name string,

Total_chats int,

Avg_Response_Time string,

Avg_Resolution_Time string,

Avg_Rating float,

Total_Feedback int

)

row format delimited

fields terminated by ','

tblproperties ("skip.header.line.count" = "1");

```
[cloudera@quickstart ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> drop database agent_details;
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask. InvalidOperationException(messag
e:Database agent_details is not empty. One or more tables exist.)
hive> drop database if exists agent_details cascade;
OK
Time taken: 16.431 seconds
hive> create database agent_details;
OK
Time taken: 1.245 seconds
hive> use agent_details;
OK
Time taken: 0.3 seconds
hive> create table agent_logging
    > (
    > s_no int,
    > agent string,
    > date date,
    > login_time string,
    > logout_time string,
    > duration string
    > )
    > row format delimited
    > fields terminated by ','
    > tblproperties("skip.header.line.count"="1");
OK
Time taken: 1.998 seconds
hive> create table agent_performances
    > (
    > s_no int,
    > date date,
    > agent_name string,
    > total_chats int,
    > average_response_time string,
    > average_resolution_time string,
    > average_rating float,
    > total_feedback int
    > )
    > row format delimited
    > fields terminated by ','
    > tblproperties("skip.header.line.count"="1");
OK
Time taken: 0.985 seconds
```

2. **Dump the data inside the hdfs in the given schema location.**

load data local inpath'Agent_Details/AgentLogingReport.csv' into table agent_logging;
load data local inpath'Agent_Details/AgentPerformance.csv' into table agent_performance;
select * from agent_logging limit 3;
select * from agent_performances limit 3;

```
hive> load data inpath 'Agent_Details/AgentLogingReport.csv' into table agent_logging;
Loading data to table agent_details.agent_logging
Table agent_details.agent_logging stats: [numFiles=1, totalSize=53908]
OK
Time taken: 3.13 seconds
hive> load data inpath 'Agent_Details/AgentPerformance.csv' into table agent_performances;
Loading data to table agent_details.agent_performances
Table agent_details.agent_performances stats: [numFiles=1, totalSize=109011]
OK
Time taken: 3.643 seconds
hive> select * from agent_logging limit 3;
OK
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10
2       Khushboo Priya  2022-07-30      15:06:59        15:07:16        0:00:17
3       Nandani Gupta   2022-07-30      15:04:24        17:31:07        2:26:42
Time taken: 4.037 seconds, Fetched: 3 row(s)
hive> select * from agent_performances limit 3;
OK
1       2022-07-30      Prerna Singh    11      0:00:38 0:04:20 4.11    9
2       2022-07-30      Nandani Gupta   11      0:01:15 0:28:25 3.14    7
3       2022-07-30      Ameya Jain      14      0:00:30 0:11:36 4.55    11
Time taken: 0.796 seconds, Fetched: 3 row(s)
```

3. **List of all agents' names.**

hive> select distinct agent from agent_logging;

```
hive> select distinct agent from agent_logging;
Query ID = cloudera_20230321051919_9c512323-2899-4fe8-96f9-139b1a4beceb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_167939269512
3_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:22:05,947 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:22:43,998 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.26 sec
2023-03-21 05:22:58,560 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.96 sec
MapReduce Total cumulative CPU time: 3 seconds 960 msec
Ended Job = job_1679392695123_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.96 sec   HDFS Read: 61634 HDFS Write: 638 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 960 msec
OK
Aditya Shinde
Aditya_iot
Amersh
Ameya Jain
Ankitjha
Anurag Tiwari
Aravind
Ayushi Mishra
Bharath
Boktiar Ahmed Bappy
Chaitra K Hiremath
Deepranjan Gupta
Dibyanshu
Harikrishnan Shaji
Hrisikesh Neogi
Hyder Abbas
Ineuron Intelligence
Ishawant Kumar
Jawala Prakash
Jaydeep Dixit
Khushboo Priya
Madhulika G
Mahesh Sarade
```

2 Items in Trash

hive> select distinct agent_name from agent_performance;

```
hive> select distinct agent_name from agent_performances;
Query ID = cloudera_20230321052626_569361eb-ced1-4b3f-8770-22e9cd3920e3
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_167939269512
3_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0002
```

```
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:26:23,061 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:26:42,479 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.16 sec
2023-03-21 05:26:59,518 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.28 sec
MapReduce Total cumulative CPU time: 4 seconds 280 msec
Ended Job = job_1679392695123_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.28 sec   HDFS Read: 117459 HDFS Write: 840 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 280 msec
OK
Abhishek
Aditya
Aditya Shinde
Aditya_iot
Amersh
Ameya Jain
Anirudh
Ankit Sharma
Ankitjha
Anurag Tiwari
Aravind
Ashad Nasim
Ashish
Ayushi Mishra
Bharath
Boktiar Ahmed Bappy
Chaitra K Hiremath
Deepranjan Gupta
Dibyanshu
Harikrishnan Shaji
Hitesh Choudhary
Hrisikesh Neogi
Hyder Abbas
Ineuron Intelligence
Ishawant Kumar
Jawala Prakash
Jayant Kumar
Jaydeep Dixit
Khushboo Priya
Madhulika G
Mahak
Mahesh Sarade
Maitry
Maneesh
Manjunatha A
Mithun S
```

4. **Find out agent average rating.**

hive> set hive.cli.print.header = true;
hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance group by agent_name;

```
hive> set hive.cli.print.header=true;
hive> select agent_name as Agents, avg(average_rating) as average_rating from agent_performances group by agent_name
;
Query ID = cloudera_20230321053131_56a2479f-31ee-4841-b415-a3639b99c030
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_167939269512
3_0003/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:31:36,233 Stage-1 map = 0%,  reduce = 0%
```

```
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:31:36,233 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:31:51,426 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.94 sec
2023-03-21 05:32:09,479 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.02 sec
MapReduce Total cumulative CPU time: 4 seconds 20 msec
Ended Job = job_1679392695123_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.02 sec   HDFS Read: 118707 HDFS Write: 1877 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 20 msec
OK
agents  average_rating
Abhishek        0.0
Aditya  0.0
Aditya Shinde   1.8003333409627278
Aditya_iot      2.3453333377838135
Amersh  0.0
Ameya Jain      2.21966667175293
Anirudh 0.6449999968210857
Ankit Sharma    0.0
Ankitjha        0.26666666666666666
Anurag Tiwari   0.18333333333333332
Aravind 2.1813333511352537
Ashad Nasim     0.16666666666666666
Ashish  0.0
Ayushi Mishra   3.481999969482422
Bharath 2.9836666584014893
Boktiar Ahmed Bappy     3.567999982833862
Chaitra K Hiremath      0.8646666606267294
Deepranjan Gupta        2.886666695276896
Dibyanshu       0.0
Harikrishnan Shaji      2.6396666526794434
Hitesh Choudhary        0.0
Hrisikesh Neogi 3.1363333304723104
Hyder Abbas     0.0
Ineuron Intelligence    0.0
Ishawant Kumar  3.543333347638448
Jawala Prakash  3.472000018755595
Jayant Kumar    1.068666664759318
Jaydeep Dixit   3.1670000314712525
Khushboo Priya  3.703666663169861
Madhulika G     3.4986666520436605
Mahak   0.1
Mahesh Sarade   2.4003333330154417
Maitry  2.9270000139872234
Maneesh 0.16666666666666666
Manjunatha A    3.5946666876475017
Mithun S        2.359000023206075
Mukesh  0.3096666653951009
```


```
Mukesh Rao      0.25566666523615517
Muskan Garg     0.712333329518636
Nandani Gupta   2.9236666679382326
Nishtha Jain    3.282333334287008
Nitin M 0.0
Prabir Kumar Satapathy  2.5103333314259846
Prateek _iot    2.4383333206176756
Prerna Singh    3.2326666434605915
Rishav Dash     1.4268333355585734
Rohan   0.0
Saif Khan       0.0
Saikumarreddy N 1.9803333441416422
Samprit 0.0
Sandipan Saha   0.4289999961853027
Sanjeev Kumar   3.3830000241597493
Sanjeevan       0.0
Saurabh Shukla  0.5556666692097981
Shiva Srivastava        0.9446666717529297
Shivan K        2.841333341598511
Shivan_S        0.14166666666666666
Shivananda Sonwane      4.232666659355163
Shubham Sharma  3.2253333568572997
Sowmiya Sivakumar       1.2599999984105428
Spuri   0.0
Sudhanshu Kumar 0.3333333333333333
Suraj S Bilgi   0.31200000445048015
Swati   2.4236666917800904
Tarun   0.05
Uday Mishra     0.0
Vasanth P       0.0
Vivek   0.5006666660308838
Wasim   2.400000015894572
Zeeshan 2.286999988555908
Time taken: 49.906 seconds, Fetched: 70 row(s)
hive>
```

5. **Total working days for each agent**

hive> select agent_name as agent, count(distinct date) as number_of_working_days from agent_performance group by agent_name;

```
hive> select agent_name as agents, count(distinct date) as no_of_working_days from agent_performances group by agent_name;
Query ID = cloudera_20230321053737_456abfb5-3aa0-4720-b7a3-aeb8e1836928
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0004, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_0004/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0004
```

```
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:37:19,451 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:37:38,109 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.36 sec
2023-03-21 05:37:53,985 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.4 sec
MapReduce Total cumulative CPU time: 4 seconds 400 msec
Ended Job = job_1679392695123_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.4 sec   HDFS Read: 118393 HDFS Write: 1050 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 400 msec
OK
Abhishek        30
Aditya  30
Aditya Shinde   30
Aditya_iot      30
Amersh  30
Ameya Jain      30
Anirudh 30
Ankit Sharma    30
Ankitjha        30
Anurag Tiwari   30
Aravind 30
Ashad Nasim     30
Ashish  30
Ayushi Mishra   30
Bharath 30
Boktiar Ahmed Bappy     30
Chaitra K Hiremath      30
Deepranjan Gupta        30
Dibyanshu       30
Harikrishnan Shaji      30
Hitesh Choudhary        30
Hrisikesh Neogi 30
Hyder Abbas     30
Ineuron Intelligence    30
Ishawant Kumar  30
Jawala Prakash  30
Jayant Kumar    30
Jaydeep Dixit   30
Khushboo Priya  30
Madhulika G     30
```

```
Madhulika G     30
Mahak   30
Mahesh Sarade   30
Maitry  30
Maneesh 30
Manjunatha A    30
Mithun S        30
Mukesh  30
Mukesh Rao      30
Muskan Garg     30
Nandani Gupta   30
Nishtha Jain    30
Nitin M 30
Prabir Kumar Satapathy  30
Prateek _iot    30
Prerna Singh    30
Rishav Dash     30
Rohan   30
Saif Khan       30
Saikumarreddy N 30
Samprit 30
Sandipan Saha   30
Sanjeev Kumar   30
Sanjeevan       30
Saurabh Shukla  30
Shiva Srivastava        30
Shivan K        30
Shivan_S        30
Shivananda Sonwane      30
Shubham Sharma  30
Sowmiya Sivakumar       30
Spuri   30
Sudhanshu Kumar 30
Suraj S Bilgi   30
Swati   30
Tarun   30
Uday Mishra     30
Vasanth P       30
Vivek   30
Wasim   30
Zeeshan 30
Time taken: 56.25 seconds, Fetched: 70 row(s)
hive>
```

6. Total query that each agent has taken

hive> select agent_name as agent, sum(total_chats) as queries_taken from agent_performance group by agent_name;

```
hive> select agent_name as agents, sum(total_chats) as total_queries_taken from agent_performances group by agent_name;
Query ID = cloudera_20230321054242_1b5603aa-fec5-4dd8-8d40-8282bd1d2f6c
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0005, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_0005/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0005
```

```
Mithun S        503
Mukesh  19
Mukesh Rao      5
Muskan Garg     56
Nandani Gupta   560
Nishtha Jain    373
Nitin M 0
Prabir Kumar Satapathy  299
Prateek _iot    190
Prerna Singh    401
Rishav Dash     409
Rohan   0
Saif Khan       0
Saikumarreddy N 364
Samprit 1
Sandipan Saha   30
Sanjeev Kumar   507
Sanjeevan       0
Saurabh Shukla  16
Shiva Srivastava        53
Shivan K        357
Shivan_S        7
Shivananda Sonwane      441
Shubham Sharma  510
Sowmiya Sivakumar       206
Spuri   0
Sudhanshu Kumar 2
Suraj S Bilgi   28
Swati   524
Tarun   22
Uday Mishra     0
Vasanth P       0
Vivek   44
Wasim   433
Zeeshan 542
Time taken: 50.069 seconds, Fetched: 70 row(s)
hive>
```

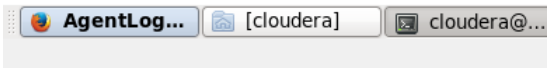

7. **Total Feedback that each agent has received**

hive> select agent_name as agent, sum(total_feedback) as feedbacks_received from agent_performance group by agent_name;

```
hive> select agent_name as agents, sum(total_feedback) as total_feedbacks from agent_performances group by agent_name;
Query ID = cloudera_20230321054747_b4945159-9723-49eb-b034-028ca774f335
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_0006/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:47:19,158 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:47:32,644 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.77 sec
2023-03-21 05:47:50,786 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.54 sec
MapReduce Total cumulative CPU time: 3 seconds 540 msec
Ended Job = job_1679392695123_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.54 sec   HDFS Read: 118229 HDFS Write: 1057 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 540 msec
OK
Abhishek        0
Aditya  0
Aditya Shinde   153
Aditya_iot      131
Amersh  0
Ameya Jain      228
Anirudh 39
Ankit Sharma    0
Ankitjha        3
Anurag Tiwari   3
Aravind 233
Ashad Nasim     9
Ashish  0
Ayushi Mishra   329
Bharath 247
Boktiar Ahmed Bappy     311
Chaitra K Hiremath      37
Deepranjan Gupta        312
Dibyanshu       0
```

AgentLog... | [cloudera] | cloudera@... | Hive-Mini-P... | [AgentPerf... | Downloads

```
Hrisikesh Neogi 367
Hyder Abbas     0
Ineuron Intelligence    0
Ishawant Kumar  202
Jawala Prakash  250
Jayant Kumar    70
Jaydeep Dixit   305
Khushboo Priya  289
Madhulika G     281
Mahak   5
Mahesh Sarade   216
Maitry  347
Maneesh 3
Manjunatha A    254
Mithun S        364
Mukesh  17
Mukesh Rao      5
Muskan Garg     37
Nandani Gupta   308
Nishtha Jain    257
Nitin M 0
Prabir Kumar Satapathy  222
Prateek _iot    107
Prerna Singh    235
Rishav Dash     264
Rohan   0
Saif Khan       0
Saikumarreddy N 290
Samprit 0
Sandipan Saha   18
Sanjeev Kumar   311
Sanjeevan       0
Saurabh Shukla  8
Shiva Srivastava        46
Shivan K        243
Shivan_S        4
Shivananda Sonwane      263
Shubham Sharma  300
Sowmiya Sivakumar       141
Spuri   0
Sudhanshu Kumar 2
Suraj S Bilgi   15
Swati   302
Tarun   6
Uday Mishra     0
Vasanth P       0
Vivek   20
Wasim   284
Zeeshan 335
Time taken: 48.228 seconds, Fetched: 70 row(s)
hive>
```

8. **Agent name who have average rating between 3.5 to 4**

hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance group by agent_name having average_rating between 3.5 and 4;

```
hive> select agent_name as agents, avg(average_rating) as average_rating from agent_performances group by agent_name h
aving average_rating between 3.5 and 4;
Query ID = cloudera_20230321055252_ffd11884-3fa2-42ce-a2d4-6242ab78ae22
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_
0007/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:52:27,946 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:52:42,944 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.76 sec
2023-03-21 05:52:59,935 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.17 sec
MapReduce Total cumulative CPU time: 4 seconds 170 msec
Ended Job = job_1679392695123_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.17 sec   HDFS Read: 119332 HDFS Write: 136 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 170 msec
OK
Boktiar Ahmed Bappy     3.567999982833862
Ishawant Kumar  3.543333347638448
Khushboo Priya  3.703666663169861
Manjunatha A    3.5946666876475017
Time taken: 48.381 seconds, Fetched: 4 row(s)
hive>
```

9. **Agent name who have rating less than 3.5**

hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance group by agent_name having average_rating < 3.5;

```
hive> select agent_name as agents, avg(average_rating) as average_rating from agent_performances group by agent_name having average_rating < 3.5;
Query ID = cloudera_20230321055555_5fc7bfe5-0454-46c8-9ef7-bf01013ac199
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_0008/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 05:55:45,309 Stage-1 map = 0%,  reduce = 0%
2023-03-21 05:55:56,648 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.54 sec
2023-03-21 05:56:11,941 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.65 sec
MapReduce Total cumulative CPU time: 3 seconds 650 msec
Ended Job = job_1679392695123_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.65 sec   HDFS Read: 119135 HDFS Write: 1704 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 650 msec
OK
Abhishek        0.0
Aditya  0.0
Aditya Shinde   1.8003333409627278
Aditya_iot      2.3453333377838135
Amersh  0.0
Ameya Jain      2.21966667175293
Anirudh 0.6449999968210857
Ankit Sharma    0.0
Ankitjha        0.26666666666666666
Anurag Tiwari   0.18333333333333332
Aravind 2.1813333511352537
Ashad Nasim     0.16666666666666666
Ashish  0.0
Ayushi Mishra   3.481999969482422
Bharath 2.9836666584014893
Chaitra K Hiremath      0.8646666606267294
Deepranjan Gupta        2.886666695276896
Dibyanshu       0.0
Harikrishnan Shaji      2.6396666526794434
Hitesh Choudhary        0.0
Hrisikesh Neogi 3.1363333304723104
Hyder Abbas     0.0
Ineuron Intelligence    0.0
Jawala Prakash  3.472000018755595
Jayant Kumar    1.068666664759318
```

```
Jayant Kumar    1.068666664759318
Jaydeep Dixit   3.1670000314712525
Madhulika G     3.4986666520436605
Mahak   0.1
Mahesh Sarade   2.4003333330154417
Maitry  2.9270000139872234
Maneesh 0.16666666666666666
Mithun S        2.359000023206075
Mukesh  0.3096666653951009
Mukesh Rao      0.25566666523615517
Muskan Garg     0.712333329518636
Nandani Gupta   2.9236666679382326
Nishtha Jain    3.282333334287008
Nitin M 0.0
Prabir Kumar Satapathy  2.5103333314259846
Prateek _iot    2.4383333206176756
Prerna Singh    3.2326666434605915
Rishav Dash     1.4268333355585734
Rohan   0.0
Saif Khan       0.0
Saikumarreddy N 1.9803333441416422
Samprit 0.0
Sandipan Saha   0.4289999961853027
Sanjeev Kumar   3.3830000241597493
Sanjeevan       0.0
Saurabh Shukla  0.5556666692097981
Shiva Srivastava        0.9446666717529297
Shivan K        2.841333341598511
Shivan_S        0.14166666666666666
Shubham Sharma  3.2253333568572997
Sowmiya Sivakumar       1.2599999984105428
Spuri   0.0
Sudhanshu Kumar 0.3333333333333333
Suraj S Bilgi   0.31200000445048015
Swati   2.4236666917800904
Tarun   0.05
Uday Mishra     0.0
Vasanth P       0.0
Vivek   0.5006666660308838
Wasim   2.400000015894572
Zeeshan 2.286999988555908
Time taken: 40.94 seconds, Fetched: 65 row(s)
hive>
```

10. **Agent name who have rating more than 4.5**

hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance group by agent_name having average_rating > 4.5;

```
hive> select agent_name as agents, avg(average_rating) as average_rating from agent_performances group by agent_name having average_rating > 4.5;
Query ID = cloudera_20230321061616_2df6830a-5695-4275-8077-3f96e0b4c8e4
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679392695123_0010, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679392695123_0010/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679392695123_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-21 06:16:58,070 Stage-1 map = 0%,  reduce = 0%
2023-03-21 06:17:11,741 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.85 sec
2023-03-21 06:17:29,576 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.93 sec
MapReduce Total cumulative CPU time: 3 seconds 930 msec
Ended Job = job_1679392695123_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.93 sec   HDFS Read: 119141 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 930 msec
OK
Time taken: 44.354 seconds
hive>
```

11. **How many feedback agents have received more than 4.5 average**

hive> select count(*) from(select agent_name as agent, avg(total_feedback) as average_feedback from agent_performance group by agent_name having average_feedback > 4.5)t;

```
hive> select agent name, count(total feedback) as total feedback from agent performances where average rating > 4.5 group by agent name order by total feedback desc;
Query ID = cloudera_20230322022727_0101a5ed-4b09-4ddc-ba2f-634924b572d5
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0007/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-22 02:27:46,119 Stage-1 map = 0%,  reduce = 0%
2023-03-22 02:27:56,481 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.75 sec
2023-03-22 02:28:07,391 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.1 sec
MapReduce Total cumulative CPU time: 3 seconds 100 msec
Ended Job = job_1679467492549_0007
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0008/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0008
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2023-03-22 02:28:19,029 Stage-2 map = 0%,  reduce = 0%
```

```
Total MapReduce CPU Time Spent: 5 seconds 810
OK
Shivananda Sonwane      17
Bharath 17
Khushboo Priya  13
Manjunatha A    13
Jaydeep Dixit   13
Ishawant Kumar  12
Shubham Sharma  12
Aravind 11
Sanjeev Kumar   11
Hrisikesh Neogi 11
Wasim   10
Prerna Singh    9
Prateek _iot    9
Shivan K        9
Ayushi Mishra   8
Ameya Jain      8
Boktiar Ahmed Bappy     8
Saikumarreddy N 8
Madhulika G     8
Swati   7
Nandani Gupta   7
Harikrishnan Shaji      7
Deepranjan Gupta        7
Aditya Shinde   7
Maitry  7
Aditya_iot      6
Nishtha Jain    6
Prabir Kumar Satapathy  6
Mithun S        5
Jawala Prakash  5
Mahesh Sarade   4
Zeeshan 3
Rishav Dash     3
Mukesh  2
Muskan Garg     2
Chaitra K Hiremath      2
Saurabh Shukla  2
Jayant Kumar    2
Sudhanshu Kumar 2
Shiva Srivastava        1
Vivek   1
Sowmiya Sivakumar       1
Mukesh Rao      1
Ankitjha        1
Sandipan Saha   1
Anirudh 1
Suraj S Bilgi   1
Time taken: 65.305 seconds, Fetched: 47 row(s)
hive>
```

12. **average weekly response time for each agent**

hive> select agent, avg(weekly_response_time_in_sec) as avg_weekly_response_time_in_sec from (select week, agent, sum((time[0]*3600+time[1]*60+time[2])) as weekly_response_time_in_sec

from(select agent_name as agent, weekofyear(date) as week, split(average_response_time,':') as time from agent_performance) t group by agent, week)s group by agent;

```
hive> select s.agent_name, avg(col1[0]*3600+col1[1]*60+substr(col1[2],1,2))/3600 from(select agent_name, split(average_resp
onse_time,':') as col1 from agent_performances)s group by s.agent_name;
Query ID = cloudera_20230322065656_412092d7-a057-4e38-8985-5b5963782fbf
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0013, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0013/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0013
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-22 06:56:38,591 Stage-1 map = 0%,  reduce = 0%
2023-03-22 06:56:49,536 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.02 sec
2023-03-22 06:57:01,653 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.94 sec
MapReduce Total cumulative CPU time: 3 seconds 940 msec
Ended Job = job_1679467492549_0013
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.94 sec   HDFS Read: 121895 HDFS Write: 1964 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 940 msec
OK
Abhishek        0.0
Aditya  0.0
Aditya Shinde   0.00825925925925926
Aditya_iot      0.009435185185185185
Amersh  0.0
Ameya Jain      0.00587037037037037
Anirudh 0.006046296296296296
Ankit Sharma    0.0
Ankitjha        0.0012314814814814816
Anurag Tiwari   0.0023425925925925927
Aravind 0.005935185185185186
Ashad Nasim     0.01073148148148148


Hrisikesh Neogi 0.014027777777777778
Hyder Abbas     0.0
Ineuron Intelligence     0.0
Ishawant Kumar  0.013925925925925925
Jawala Prakash  0.026175925925925925
Jayant Kumar    0.005120370370370371
Jaydeep Dixit   0.012333333333333333
Khushboo Priya  0.017027777777777777
Madhulika G     0.018453703703703705
Mahak    0.0
Mahesh Sarade   0.012898148148148146
Maitry  0.017731481481481483
Maneesh 0.00125
Manjunatha A    0.010046296296296296
Mithun S        0.008037037037037037
Mukesh  9.25925925925926E-4
Mukesh Rao      0.003648148148148148
Muskan Garg     0.0016481481481481482
Nandani Gupta   0.01662962962962963
Nishtha Jain    0.016888888888888887
Nitin M 0.0
Prabir Kumar Satapathy  0.010555555555555556
Prateek _iot    0.00625
Prerna Singh    0.01324074074074074
Rishav Dash     0.008421296296296297
Rohan    0.0
Saif Khan       0.0
Saikumarreddy N 0.006990740740740741
Samprit 0.0
Sandipan Saha   0.001638888888888889
Sanjeev Kumar   0.014222222222222223
Sanjeevan       0.0
Saurabh Shukla  9.7222222222222E-4
Shiva Srivastava        0.002777777777777778
Shivan K        0.013305555555555555
Shivan_S        6.759259259259258E-4
Shivananda Sonwane      0.015555555555555555
Shubham Sharma  0.013425925925925926
Sowmiya Sivakumar       0.007268518518518519
Spuri    0.0
Sudhanshu Kumar 0.0011111111111111111
Suraj S Bilgi   0.0016851851851851852
Swati    0.016055555555555556
Tarun    0.0
Uday Mishra     0.0
Vasanth P       0.0
Vivek    0.0038055555555555555
Wasim    0.00825
Zeeshan 0.01714814814814815
Time taken: 35.347 seconds, Fetched: 70 row(s)
hive>
```

13. **average weekly resolution time for each agents**

hive> select agent, avg(weekly_resolution_time_in_sec) as avg_weekly_resolution_time_in_sec from (select week, agent, sum((time[0]*3600+time[1]*60+time[2])) as weekly_resolution_time_in_sec from(select agent_name as agent, weekofyear(date) as week, split(average_resolution_time,':') as time from agent_performance) t group by agent, week)s group by agent;

```
hive> select s.agent_name, avg(col1[0]*3600+col1[1]*60+substr(col1[2],1,2))/3600 from(select agent_name, split(average_reso
lution_time,':') as col1 from agent_performances)s group by s.agent_name;
Query ID = cloudera_20230322072424_d9393971-dbad-424d-aafe-f0387c1f2312
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0021, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0021/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0021
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-22 07:24:51,829 Stage-1 map = 0%,  reduce = 0%
2023-03-22 07:25:03,375 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.09 sec
2023-03-22 07:25:15,388 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.18 sec
MapReduce Total cumulative CPU time: 4 seconds 180 msec
Ended Job = job_1679467492549_0021
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.18 sec   HDFS Read: 121895 HDFS Write: 1995 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 180 msec
OK
Abhishek        0.0
Aditya  0.0
Aditya Shinde   0.17239814814814813
Aditya_iot      0.16369444444444442
Amersh  0.0
Ameya Jain      0.09125
Anirudh 0.05151851851851852
Ankit Sharma    0.0
Ankitjha        0.01512962962962963
Anurag Tiwari   0.020537037037037038
Aravind 0.14876851851851852
Ashad Nasim     0.005814814814814815
Ashish  0.0
Ayushi Mishra   0.25627777777777777
Bharath 0.1798888888888889
Boktiar Ahmed Bappy     0.2843981481481482
Chaitra K Hiremath      0.024712962962962964
Deepranjan Gupta        0.3475462962962963
Dibyanshu       0.006851851851851852
Harikrishnan Shaji      0.18845370370370368
Hitesh Choudhary        7.87037037037037E-4
Hrisikesh Neogi 0.25730555555555557
```

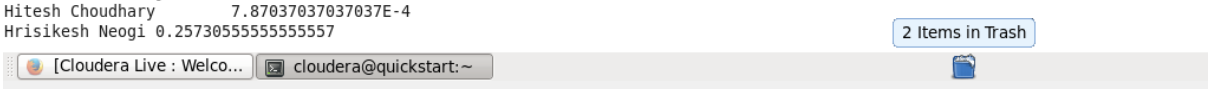

14. **Find the number of chats on which they have received a feedback**

hive> select agent_name as agent, sum(total_chats) as chats, sum(total_feedback) as chats_with_feedback_received from agent_performance group by agent_name;

```
hive> select agent_name, sum(total_chats) as chats,sum(total_feedback) as total_feedbacks from agent_performances group by
agent_name;
Query ID = cloudera_20230322092727_1adc00b3-497f-4d7a-a87a-6c0ffd31bc24
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0029, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0029/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0029
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-22 09:27:12,799 Stage-1 map = 0%,  reduce = 0%
2023-03-22 09:27:24,326 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.39 sec
2023-03-22 09:27:37,450 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 2.92 sec
MapReduce Total cumulative CPU time: 2 seconds 920 msec
Ended Job = job_1679467492549_0029
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 2.92 sec   HDFS Read: 118676 HDFS Write: 1278 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 920 msec
OK
agent_name      chats   total_feedbacks
Abhishek        0       0
Aditya  0       0
Aditya Shinde   277     153
Aditya_iot      231     131
Amersh  0       0
Ameya Jain      322     228
Anirudh 81      39
Ankit Sharma    0       0
Ankitjha        5       3
Anurag Tiwari   4       3
Aravind 366     233
Ashad Nasim     18      9
Ashish  0       0
Ayushi Mishra   514     329
Bharath 369     247
Boktiar Ahmed Bappy     452     311
Chaitra K Hiremath      64      37
Deepranjan Gupta        493     312
Dibyanshu       1       0
Harikrishnan Shaji      381     231
Hitesh Choudhary        1       0
Hrisikesh Neogi 578     367
Hyder Abbas     0       0
Ineuron Intelligence    0       0
Ishawant Kumar  338     202
```

[Cloudera Live : Welco...] [cloudera@quickstart:~]

15. **Total contribution hour for each and every agents weekly basis**

hive> select week, agent, sum((time[0]*3600+time[1]*60+time[2])/3600) as total_hrs_contributed from(select agent, weekofyear(date) as week, split(duration,':') as time from agent_loging) t group by agent, week;

```
hive> select s.agent, sum(col1[0]*3600+col1[1]*60+col1[2])/3600 total_hours, s.weekly from (select agent, split(duration,':
') as col1, weekofyear(date) as weekly from agent_logging)s group by s.agent, s.weekly;
Query ID = cloudera_20230322092424_ccb05267-9c2b-4a1e-b60a-6914c23ec675
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1679467492549_0028, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0028/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0028
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-03-22 09:24:15,641 Stage-1 map = 0%,  reduce = 0%
2023-03-22 09:24:26,959 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.99 sec
2023-03-22 09:24:38,991 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 3.99 sec
MapReduce Total cumulative CPU time: 3 seconds 990 msec
Ended Job = job_1679467492549_0028
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.99 sec   HDFS Read: 66162 HDFS Write: 3035 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 990 msec
OK
s.agent total_hours     s.weekly
Aditya Shinde   0.03611111111111111     30
Aditya_iot      6.095277777777778       29
Aditya_iot      9.635833333333334       30
Amersh  3.063888888888889       30
Ameya Jain      24.083055555555557      29
Ameya Jain      17.9925 30
Ankitjha        2.2669444444444444      30
Anurag Tiwari   0.2644444444444444      29
Anurag Tiwari   2.5144444444444445      30
Aravind 24.235555555555557      29
Aravind 0.0636111111111111      30
Ayushi Mishra   17.790277777777778      29
Ayushi Mishra   20.33138888888889       30
Bharath 24.070833333333333      29
Bharath 24.00583333333333       30
Boktiar Ahmed Bappy     17.75027777777778       29
Boktiar Ahmed Bappy     22.518333333333334      30
Chaitra K Hiremath      2.234722222222222       29
Chaitra K Hiremath      32.090833333333336      30
Deepranjan Gupta        48.99638888888889       29
Deepranjan Gupta        57.278888888888886      30
Dibyanshu       27.74388888888889       29
```

16. **Perform inner join, left join and right join based on the agent column and after joining the table export that data into your local system.**

    Inner Join:

    hive> select l.*, p.* > from > agent_loging l > inner join > agent_performance p > on l.agent = p.agent_name > limit 3;

```
hive> select l.*,p.* from agent_logging l inner join agent_performances p on l.agent=p.agent_name limit 3;
Query ID = cloudera_20230322093030_3158a588-0f21-4ad2-b14c-e9695885299b
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322093030_3158a588-0f21-4ad2-b14c-e9695885299b.log
2023-03-22 09:30:23     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 09:30:26     Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/a72b7dc9-5850-403
7-8732-74c59d57f0fd/hive_2023-03-22_09-30-15_652_7987760937020627659-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.h
ashtable
2023-03-22 09:30:26     Uploaded 1 File to: file:/tmp/cloudera/a72b7dc9-5850-4037-8732-74c59d57f0fd/hive_2023-03-22_09-30-1
5_652_7987760937020627659-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (37895 bytes)
2023-03-22 09:30:26     End of local task; Time Taken: 2.519 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0030, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0030/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0030
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 09:30:39,478 Stage-3 map = 0%,  reduce = 0%
2023-03-22 09:30:49,378 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.91 sec
MapReduce Total cumulative CPU time: 1 seconds 910 msec
Ended Job = job_1679467492549_0030
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.91 sec   HDFS Read: 13371 HDFS Write: 315 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 910 msec
OK
l.s_no  l.agent l.date  l.login_time    l.logout_time   l.duration      p.s_no  p.date  p.agent_name    p.total_chats   p.a
verage_response_time    p.average_resolution_time       p.average_rating        p.total_feedback
16      Prerna Singh    2022-07-30      12:32:28        14:10:08        1:37:40 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
75      Prerna Singh    2022-07-29      17:47:06        21:03:44        3:16:37 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
91      Prerna Singh    2022-07-29      15:08:22        17:20:49        2:12:27 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
Time taken: 35.862 seconds, Fetched: 3 row(s)
hive>
```

[Cloudera Live : Welco...] [cloudera@quickstart:~]

Left Join:

hive> select l.*, p.* > from > agent_loging l > left join > agent_performance p > on l.agent = p.agent_name > limit 3;

```
hive> select l.*,p.* from agent_logging l left join agent_performances p on l.agent=p.agent_name limit 3;
Query ID = cloudera_20230322093232_4736ad4f-85b5-4acd-aa43-229bc54fa1b6
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322093232_4736ad4f-85b5-4acd-aa43-229bc54fa1b6.log
2023-03-22 09:32:50     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 09:32:52     Dump the side-table for tag: 1 with group count: 70 into file: file:/tmp/cloudera/a72b7dc9-5850-403
7-8732-74c59d57f0fd/hive_2023-03-22_09-32-42_428_8442372344074333733-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.h
ashtable
2023-03-22 09:32:52     Uploaded 1 File to: file:/tmp/cloudera/a72b7dc9-5850-4037-8732-74c59d57f0fd/hive_2023-03-22_09-32-4
2_428_8442372344074333733-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.hashtable (73261 bytes)
2023-03-22 09:32:52     End of local task; Time Taken: 2.48 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0031, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0031/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0031
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 09:33:06,489 Stage-3 map = 0%,  reduce = 0%
2023-03-22 09:33:15,746 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.64 sec
MapReduce Total cumulative CPU time: 1 seconds 640 msec
Ended Job = job_1679467492549_0031
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.64 sec   HDFS Read: 13232 HDFS Write: 348 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 640 msec
OK
l.s_no  l.agent l.date  l.login_time    l.logout_time   l.duration      p.s_no  p.date  p.agent_name    p.total_chats   p.a
verage_response_time    p.average_resolution_time       p.average_rating        p.total_feedback
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 69      2022-07-30      Shivananda
Sonwane 4       0:01:14 0:16:53 5.0     1
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 73      2022-07-29      Shivananda
Sonwane 14      0:00:45 0:15:38 4.67    9
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 214     2022-07-28      Shivananda
Sonwane 5       0:00:31 0:38:04 5.0     4
Time taken: 35.469 seconds, Fetched: 3 row(s)
hive>
```

[Cloudera Live : Welco...] [cloudera@quickstart:~]

Right Join:

hive> select l.*, p.* > from > agent_loging l > right join > agent_performance p > on l.agent = p.agent_name > limit 3;

```
hive> select l.*,p.* from agent_logging l right join agent_performances p on l.agent=p.agent_name limit 3;
Query ID = cloudera_20230322093333_0df41b0a-0d94-4387-9a33-3f86bc12bee1
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322093333_0df41b0a-0d94-4387-9a33-3f86bc12bee1.log
2023-03-22 09:34:07     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 09:34:09     Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/a72b7dc9-5850-403
7-8732-74c59d57f0fd/hive_2023-03-22_09-33-59_794_6377935058102073276-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile20--.h
ashtable
2023-03-22 09:34:09     Uploaded 1 File to: file:/tmp/cloudera/a72b7dc9-5850-4037-8732-74c59d57f0fd/hive_2023-03-22_09-33-5
9_794_6377935058102073276-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile20--.hashtable (37895 bytes)
2023-03-22 09:34:09     End of local task; Time Taken: 2.111 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0032, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0032/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0032
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 09:34:22,440 Stage-3 map = 0%,  reduce = 0%
2023-03-22 09:34:31,607 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.62 sec
MapReduce Total cumulative CPU time: 1 seconds 620 msec
Ended Job = job_1679467492549_0032
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.62 sec   HDFS Read: 13206 HDFS Write: 315 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 620 msec
OK
l.s_no  l.agent l.date  l.login_time    l.logout_time   l.duration      p.s_no  p.date  p.agent_name    p.total_chats   p.a
verage_response_time    p.average_resolution_time       p.average_rating        p.total_feedback
16      Prerna Singh    2022-07-30      12:32:28        14:10:08        1:37:40 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
75      Prerna Singh    2022-07-29      17:47:06        21:03:44        3:16:37 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
91      Prerna Singh    2022-07-29      15:08:22        17:20:49        2:12:27 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
Time taken: 33.967 seconds, Fetched: 3 row(s)
hive>
```

[Cloudera Live : Welco...] [cloudera@quickstart:~]

Export data into local system:

Inner Join:

[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_loging l inner join agent.agent_performance p on l.agent = p.agent_name limit 10' > /home/cloudera/Hive_Mini_Proj-1/agent_inner_join.csv

```
[cloudera@quickstart ~]$ hive -e 'select l.*,p.* from agent_details.agent_logging l inner join agent_details.agent_performa
nces p on l.agent=p.agent_name limit 10' > /home/cloudera/Hive-Mini-Proj-1/agent_inner_join.csv

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
Query ID = cloudera_20230322100404_adde67fb-1de3-4526-9203-f15c221fca04
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322100404_adde67fb-1de3-4526-9203-f15c221fca04.log
2023-03-22 10:04:26     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 10:04:28     Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/94675fc3-073f-4a2
c-a93a-0e2c90b0fa5f/hive_2023-03-22_10-04-16_466_7061998434702171190-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.h
ashtable
2023-03-22 10:04:29     Uploaded 1 File to: file:/tmp/cloudera/94675fc3-073f-4a2c-a93a-0e2c90b0fa5f/hive_2023-03-22_10-04-1
6_466_7061998434702171190-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (37895 bytes)
2023-03-22 10:04:29     End of local task; Time Taken: 2.738 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0033, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0033/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0033
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 10:04:44,286 Stage-3 map = 0%,  reduce = 0%
2023-03-22 10:04:54,922 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.91 sec
MapReduce Total cumulative CPU time: 1 seconds 910 msec
Ended Job = job_1679467492549_0033
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.91 sec   HDFS Read: 13199 HDFS Write: 1057 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 910 msec
OK
Time taken: 39.659 seconds, Fetched: 10 row(s)
```

Left Join:

[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_loging l left join agent.agent_performance p on l.agent = p.agent_name limit 10' > /home/cloudera/Hive_Mini_Proj-1/agent_left_join.csv

```
[cloudera@quickstart ~]$ hive -e 'select l.*,p.* from agent_details.agent_logging l left join agent_details.agent_performan
ces p on l.agent=p.agent_name limit 10' > /home/cloudera/Hive-Mini-Proj-1/agent_left_join.csv

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
Query ID = cloudera_20230322100808_aed251eb-9c58-4abc-a6d1-86eaa012af3c
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322100808_aed251eb-9c58-4abc-a6d1-86eaa012af3c.log
2023-03-22 10:08:31     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 10:08:34     Dump the side-table for tag: 1 with group count: 70 into file: file:/tmp/cloudera/fda467eb-1868-43f
6-84d8-acb16441e481/hive_2023-03-22_10-08-22_365_3563291185952597523-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.h
ashtable
2023-03-22 10:08:34     Uploaded 1 File to: file:/tmp/cloudera/fda467eb-1868-43f6-84d8-acb16441e481/hive_2023-03-22_10-08-2
2_365_3563291185952597523-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable (73261 bytes)
2023-03-22 10:08:34     End of local task; Time Taken: 2.503 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0034, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0034/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0034
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 10:08:47,877 Stage-3 map = 0%,  reduce = 0%
2023-03-22 10:08:58,279 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.71 sec
MapReduce Total cumulative CPU time: 1 seconds 710 msec
Ended Job = job_1679467492549_0034
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.71 sec   HDFS Read: 13060 HDFS Write: 1173 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 710 msec
OK
Time taken: 37.265 seconds, Fetched: 10 row(s)
```

Right Join:

[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_loging l right join agent.agent_performance p on l.agent = p.agent_name limit 10' > /tmp/Agent_data/agent_right_join.csv

```
[cloudera@quickstart ~]$ hive -e 'select l.*,p.* from agent_details.agent_logging l right join agent_details.agent_performa
nces p on l.agent=p.agent_name limit 10' > /home/cloudera/Hive-Mini-Proj-1/agent_right_join.csv

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
Query ID = cloudera_20230322100909_476e3b0f-32a4-488c-80bb-673f26a694c7
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20230322100909_476e3b0f-32a4-488c-80bb-673f26a694c7.log
2023-03-22 10:10:05     Starting to launch local task to process map join;      maximum memory = 1013645312
2023-03-22 10:10:07     Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/69e5e43b-81f9-4a2
5-8de6-76fe89672e3c/hive_2023-03-22_10-09-56_353_8276029881292291054-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.h
ashtable
2023-03-22 10:10:07     Uploaded 1 File to: file:/tmp/cloudera/69e5e43b-81f9-4a25-8de6-76fe89672e3c/hive_2023-03-22_10-09-5
6_353_8276029881292291054-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (37895 bytes)
2023-03-22 10:10:07     End of local task; Time Taken: 2.033 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0035, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0035/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0035
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-03-22 10:10:22,394 Stage-3 map = 0%,  reduce = 0%
2023-03-22 10:10:32,751 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 1.55 sec
MapReduce Total cumulative CPU time: 1 seconds 550 msec
Ended Job = job_1679467492549_0035
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 1.55 sec   HDFS Read: 13034 HDFS Write: 1057 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 550 msec
OK
Time taken: 37.72 seconds, Fetched: 10 row(s)
```

Check the exported data in local system:

```
[cloudera@quickstart ~]$ ls /home/cloudera/Hive-Mini-Proj-1
agent_inner_join.csv  agent_left_join.csv  AgentLogingReport.csv  AgentPerformance.csv  agent_right_join.csv  Assign-1

[cloudera@quickstart ~]$ head -3 /home/cloudera/Hive-Mini-Proj-1/agent_inner_join.csv
16      Prerna Singh    2022-07-30      12:32:28        14:10:08        1:37:40 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
75      Prerna Singh    2022-07-29      17:47:06        21:03:44        3:16:37 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
91      Prerna Singh    2022-07-29      15:08:22        17:20:49        2:12:27 1       2022-07-30      Prerna Singh    110
:00:38  0:04:20 4.11    9
[cloudera@quickstart ~]$ head -3 /home/cloudera/Hive-Mini-Proj-1/agent_left_join.csv
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 69      2022-07-30      Shivananda Sonwane      4       0:01:14 0:16:53 5.0     1
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 73      2022-07-29      Shivananda Sonwane      14      0:00:45 0:15:38 4.67    9
1       Shivananda Sonwane      2022-07-30      15:35:29        17:39:39        2:04:10 214     2022-07-28      Shivananda Sonwane      5       0:00:31 0:38:04 5.0     4
[cloudera@quickstart ~]$ head -3 /home/cloudera/Hive-Mini-Proj-1/agent_right_join.csv
16      Prerna Singh    2022-07-30      12:32:28        14:10:08        1:37:40 1       2022-07-30      Prerna Singh    11      0:00:38 0:04:20 4.11    9
75      Prerna Singh    2022-07-29      17:47:06        21:03:44        3:16:37 1       2022-07-30      Prerna Singh    11      0:00:38 0:04:20 4.11    9
91      Prerna Singh    2022-07-29      15:08:22        17:20:49        2:12:27 1       2022-07-30      Prerna Singh    11      0:00:38 0:04:20 4.11    9
[cloudera@quickstart ~]$
                                                                                                        Hive-Mini-Proj-1
[Cloudera Live : Welc...  cloudera@quickstart:~  [cloudera]  [Hive-Mini-Proj-1]
```

17. **Perform partitioning on top of the agent column and then on top of that perform bucketing for each partitioning.**

First set the below mentioned properties to be true.

```
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.patition.mode=nonstrict;
```

Create Partition_bucketed table:

```
hive> create table partition_bucketed_logging
> (
> s_no int,
> date date,
> login_time string,
> logout_time string,
> duration string
> )
> partitioned by (agent string)
> clustered by(s_no)
> into 4 buckets
> row format delimited
> fields terminated by ','
> stored as textfile;
```

Load data into Partition_bucketed table:

hive> insert into table partition_bucketed_loging partition(agent) select s_no, date, login_time, logout_time, duration, agent from agent_logging;

```
hive> use agent_details;
OK
Time taken: 0.019 seconds
hive> set hive.exec.dynamic.partition=true;
hive> [cloudera@quickstart ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> use agent_details;
OK
Time taken: 0.511 seconds
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> create table partition_bucketed_logging
    > (
    > s_no int,
    > date date,
    > login_time string,
    > logout_time string,
    > duration string
    > )
    > partitioned by (agent string)
    > clustered by (s_no)
    > into 4 buckets
    > row format delimited
    > fields terminated by ','
    >
    > stored as textfile;
OK
Time taken: 0.401 seconds
hive>  insert into table partition_bucketed_logging partition(agent)select s_no,date,login_time,logout_time,duration,agent from agent_logging;
Query ID = cloudera_20230322105353_532f44ad-5d3d-4a50-a3b7-ee1ac2fff544
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1679467492549_0036, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1679467492549_0036/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1679467492549_0036
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2023-03-22 10:53:49,485 Stage-1 map = 0%,  reduce = 0%
2023-03-22 10:54:03,201 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.56 sec
MapReduce Total cumulative CPU time: 2 seconds 560 msec
Ended Job = job_1679467492549_0036
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
```


```
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/agent_details.db/partition_bucketed_logging/.hive-staging_hive_2023-03-22_10-53-34_241_1839891603046782836-1/-ext-10000
Loading data to table agent_details.partition_bucketed_logging partition (agent=null)
         Time taken for load dynamic partitions : 8895
        Loading partition {agent=Sanjeev Kumar}
        Loading partition {agent=Aditya_iot}
        Loading partition {agent=Zeeshan}
        Loading partition {agent=Boktiar Ahmed Bappy}
        Loading partition {agent=Khushboo Priya}
        Loading partition {agent=Mahesh Sarade}
        Loading partition {agent=Shivananda Sonwane}
        Loading partition {agent=Sudhanshu Kumar}
        Loading partition {agent=Ankitjha}
        Loading partition {agent=Nishtha Jain}
        Loading partition {agent=Rishav Dash}
        Loading partition {agent=Anurag Tiwari}
        Loading partition {agent=Deepranjan Gupta}
        Loading partition {agent=Shivan K}
        Loading partition {agent=Hrisikesh Neogi}
        Loading partition {agent=Madhulika G}
        Loading partition {agent=Mithun S}
        Loading partition {agent=Nitin M}
        Loading partition {agent=Muskan Garg}
        Loading partition {agent=Swati}
        Loading partition {agent=Tarun}
        Loading partition {agent=Saurabh Shukla}
        Loading partition {agent=Prabir Kumar Satapathy}
        Loading partition {agent=Suraj S Bilgi}
        Loading partition {agent=Prerna Singh}
        Loading partition {agent=Aravind}
        Loading partition {agent=Sowmiya Sivakumar}
        Loading partition {agent=Wasim}
        Loading partition {agent=Nandani Gupta}
        Loading partition {agent=Prateek _iot}
        Loading partition {agent=Jaydeep Dixit}
        Loading partition {agent=Aditya Shinde}
        Loading partition {agent=Dibyanshu}
        Loading partition {agent=Manjunatha A}
        Loading partition {agent=Hyder Abbas}
        Loading partition {agent=Chaitra K Hiremath}
        Loading partition {agent=Ishawant Kumar}
        Loading partition {agent=Jawala Prakash}
        Loading partition {agent=Ayushi Mishra}
        Loading partition {agent=Mukesh}
        Loading partition {agent=Ineuron Intelligence}
        Loading partition {agent=Shubham Sharma}
        Loading partition {agent=Harikrishnan Shaji}
        Loading partition {agent=Saikumarreddy N}
        Loading partition {agent=Ameya Jain}
```

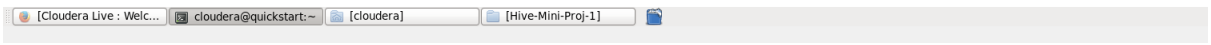

```
Partition agent_details.partition_bucketed_logging{agent=Ankitjha} stats: [numFiles=1, numRows=4, totalSize=164, rawDataSize=160]
Partition agent_details.partition_bucketed_logging{agent=Anurag Tiwari} stats: [numFiles=1, numRows=37, totalSize=1500, rawDataSize=1463]
Partition agent_details.partition_bucketed_logging{agent=Aravind} stats: [numFiles=1, numRows=10, totalSize=409, rawDataSize=399]
Partition agent_details.partition_bucketed_logging{agent=Ayushi Mishra} stats: [numFiles=1, numRows=18, totalSize=729, rawDataSize=711]
Partition agent_details.partition_bucketed_logging{agent=Bharath} stats: [numFiles=1, numRows=9, totalSize=359, rawDataSize=350]
Partition agent_details.partition_bucketed_logging{agent=Boktiar Ahmed Bappy} stats: [numFiles=1, numRows=17, totalSize=680, rawDataSize=663]
Partition agent_details.partition_bucketed_logging{agent=Chaitra K Hiremath} stats: [numFiles=1, numRows=13, totalSize=530, rawDataSize=517]
Partition agent_details.partition_bucketed_logging{agent=Deepranjan Gupta} stats: [numFiles=1, numRows=58, totalSize=2359, rawDataSize=2301]
Partition agent_details.partition_bucketed_logging{agent=Dibyanshu} stats: [numFiles=1, numRows=208, totalSize=8361, rawDataSize=8153]
Partition agent_details.partition_bucketed_logging{agent=Harikrishnan Shaji} stats: [numFiles=1, numRows=23, totalSize=931, rawDataSize=908]
Partition agent_details.partition_bucketed_logging{agent=Hrisikesh Neogi} stats: [numFiles=1, numRows=37, totalSize=1493, rawDataSize=1456]
Partition agent_details.partition_bucketed_logging{agent=Hyder Abbas} stats: [numFiles=1, numRows=2, totalSize=82, rawDataSize=80]
Partition agent_details.partition_bucketed_logging{agent=Ineuron Intelligence} stats: [numFiles=1, numRows=1, totalSize=41, rawDataSize=40]
Partition agent_details.partition_bucketed_logging{agent=Ishawant Kumar} stats: [numFiles=1, numRows=49, totalSize=1980, rawDataSize=1931]
Partition agent_details.partition_bucketed_logging{agent=Jawala Prakash} stats: [numFiles=1, numRows=16, totalSize=636, rawDataSize=620]
Partition agent_details.partition_bucketed_logging{agent=Jaydeep Dixit} stats: [numFiles=1, numRows=11, totalSize=449, rawDataSize=438]
Partition agent_details.partition_bucketed_logging{agent=Khushboo Priya} stats: [numFiles=1, numRows=18, totalSize=726, rawDataSize=708]
Partition agent_details.partition_bucketed_logging{agent=Madhulika G} stats: [numFiles=1, numRows=17, totalSize=687, rawDataSize=670]
Partition agent_details.partition_bucketed_logging{agent=Mahesh Sarade} stats: [numFiles=1, numRows=36, totalSize=1468, rawDataSize=1432]
Partition agent_details.partition_bucketed_logging{agent=Maitry} stats: [numFiles=1, numRows=5, totalSize=200, rawDataSize=195]
Partition agent_details.partition_bucketed_logging{agent=Manjunatha A} stats: [numFiles=1, numRows=8, totalSize=318, rawDataSize=310]
Partition agent_details.partition_bucketed_logging{agent=Mithun S} stats: [numFiles=1, numRows=14, totalSize=564, rawDataSize=550]
Partition agent_details.partition_bucketed_logging{agent=Mukesh} stats: [numFiles=1, numRows=3, totalSize=121, rawDataSize=118]
Partition agent_details.partition_bucketed_logging{agent=Muskan Garg} stats: [numFiles=1, numRows=12, totalSize=491, rawDataSize=479]
Partition agent_details.partition_bucketed_logging{agent=Nandani Gupta} stats: [numFiles=1, numRows=11, totalSize=447, rawDataSize=436]
Partition agent_details.partition_bucketed_logging{agent=Nishtha Jain} stats: [numFiles=1, numRows=18, totalSize=728, rawDataSize=710]
Partition agent_details.partition_bucketed_logging{agent=Nitin M} stats: [numFiles=1, numRows=1, totalSize=41, rawDataSize=40]
Partition agent_details.partition_bucketed_logging{agent=Prabir Kumar Satapathy} stats: [numFiles=1, numRows=26, totalSize=1054, rawDataSize=1028]
Partition agent_details.partition_bucketed_logging{agent=Prateek _iot} stats: [numFiles=1, numRows=17, totalSize=690, rawDataSize=673]
Partition agent_details.partition_bucketed_logging{agent=Prerna Singh} stats: [numFiles=1, numRows=18, totalSize=735, rawDataSize=717]
Partition agent_details.partition_bucketed_logging{agent=Rishav Dash} stats: [numFiles=1, numRows=12, totalSize=481, rawDataSize=469]
Partition agent_details.partition_bucketed_logging{agent=Saikumarreddy N} stats: [numFiles=1, numRows=10, totalSize=410, rawDataSize=400]
Partition agent_details.partition_bucketed_logging{agent=Sanjeev Kumar} stats: [numFiles=1, numRows=20, totalSize=815, rawDataSize=795]
Partition agent_details.partition_bucketed_logging{agent=Saurabh Shukla} stats: [numFiles=1, numRows=40, totalSize=1608, rawDataSize=1568]
Partition agent_details.partition_bucketed_logging{agent=Shiva Srivastava} stats: [numFiles=1, numRows=15, totalSize=612, rawDataSize=597]
Partition agent_details.partition_bucketed_logging{agent=Shivan K} stats: [numFiles=1, numRows=36, totalSize=1458, rawDataSize=1422]
Partition agent_details.partition_bucketed_logging{agent=Shivananda Sonwane} stats: [numFiles=1, numRows=15, totalSize=600, rawDataSize=585]
Partition agent_details.partition_bucketed_logging{agent=Shubham Sharma} stats: [numFiles=1, numRows=35, totalSize=1425, rawDataSize=1390]
Partition agent_details.partition_bucketed_logging{agent=Sowmiya Sivakumar} stats: [numFiles=1, numRows=24, totalSize=972, rawDataSize=948]
Partition agent_details.partition_bucketed_logging{agent=Sudhanshu Kumar} stats: [numFiles=1, numRows=11, totalSize=446, rawDataSize=435]
Partition agent_details.partition_bucketed_logging{agent=Suraj S Bilgi} stats: [numFiles=1, numRows=5, totalSize=200, rawDataSize=195]
Partition agent_details.partition_bucketed_logging{agent=Swati} stats: [numFiles=1, numRows=5, totalSize=205, rawDataSize=200]
Partition agent_details.partition_bucketed_logging{agent=Tarun} stats: [numFiles=1, numRows=1, totalSize=43, rawDataSize=42]
Partition agent_details.partition_bucketed_logging{agent=Wasim} stats: [numFiles=1, numRows=20, totalSize=813, rawDataSize=793]
Partition agent_details.partition_bucketed_logging{agent=Zeeshan} stats: [numFiles=1, numRows=10, totalSize=409, rawDataSize=399]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 2.56 sec   HDFS Read: 58915 HDFS Write: 44385 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 560 msec
OK
Time taken: 45.293 seconds
hive>
```